

The AutoTokenizer used with LLaMA models typically relies on **SentencePiece** tokenization, specifically using **Byte-Pair Encoding (BPE)** or **Unigram Language Model** encoding depending on the version.

The AutoTokenizer in Hugging Face Transformers automatically loads the correct tokenizer configuration based on the model you specify (e.g., "meta-llama/Llama-2-7b-hf"), and under the hood, it uses the LlamaTokenizer or LlamaTokenizerFast, both based on SentencePiece.

### **Key Technique:**

**SentencePiece + BPE (for LLaMA 1/2) or SentencePiece + Unigram (for LLaMA 3)**

### **Why SentencePiece?**

- It allows subword-level tokenization.
- Works directly on raw Unicode text (no preprocessing like whitespace splitting required).
- Ideal for multilingual models or models trained on varied data.

**complete dry run** on a more complex sentence, step-by-step, using LLaMA's AutoTokenizer (e.g., from "meta-llama/Llama-2-7b-hf").

### **Sentence:**

"The quick brown fox jumps over the lazy dog, effortlessly."

### **Step 1: Load Tokenizer**

```
from transformers import AutoTokenizer
```

```
tokenizer = AutoTokenizer.from_pretrained("meta-llama/Llama-2-7b-hf")
```

### **Step 2: Tokenize Sentence**

```
tokens = tokenizer.tokenize("The quick brown fox jumps over the lazy dog, effortlessly.")  
print(tokens)
```

## Output:

```
['_The', '_quick', '_brown', '_fox', '_jumps', '_over', '_the', '_lazy', '_dog', ',', '_effort',  
'lessly', '.']
```

## Step 3: Tokenization Logic (Dry Run):

### Pre-tokenization (with special marker):

The sentence is first transformed into:

```
"_The_quick_brown_fox_jumps_over_the_lazy_dog,_effortlessly."
```

Notice:

- `_` marks the beginning of a word.
- `SentencePiece` handles this, so words are not split by spaces directly.

### Subword Matching (via BPE):

- Each part is matched greedily to longest vocab tokens:
  - `'_The'` → one token
  - `'_quick'` → one token
  - `'_effortlessly'` → gets broken into: `'_effort'`, `'lessly'` (since `'effortlessly'` is not in vocab but parts of it are)

## Step 4: Token IDs

```
ids = tokenizer.convert_tokens_to_ids(tokens)  
print(ids)
```

### Example Output:

```
[1332, 2398, 4149, 2317, 3664, 2934, 2781, 4009, 1443, 29892, 13204, 22765, 29889]
```

(Note: IDs may vary slightly by tokenizer version.)

### Step 5: Decode Back

```
decoded = tokenizer.decode(ids)
print(decoded)
```

#### Output:

"The quick brown fox jumps over the lazy dog, effortlessly."

It reconstructs the sentence accurately.

#### Summary Table:

| Word           | Token(s)            | Token ID(s)  |
|----------------|---------------------|--------------|
| "The"          | '_The'              | 1332         |
| "quick"        | '_quick'            | 2398         |
| "brown"        | '_brown'            | 4149         |
| "fox"          | '_fox'              | 2317         |
| "jumps"        | '_jumps'            | 3664         |
| "over"         | '_over'             | 2934         |
| "the"          | '_the'              | 2781         |
| "lazy"         | '_lazy'             | 4009         |
| "dog"          | '_dog'              | 1443         |
| ","            | ','                 | 29892        |
| "effortlessly" | '_effort', 'lessly' | 13204, 22765 |
| ". "           | ','                 | 29889        |